

Champion-Model.R

gavin-kunish

2025-08-18

```
# Champion Model Plots

# ===== Libraries =====
library(tseries)

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

library(vars)

## Loading required package: MASS

## Loading required package: strucchange

## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric

## Loading required package: sandwich

## Loading required package: urca

## Loading required package: lmtest

library(forecast)
library(zoo)
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following object is masked from 'package:MASS':
##
##   select

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(scales)

# ===== Load & basic transforms (assumes monthly dates on the 1st) =====
# setwd("~/Desktop/")
data <- read.csv("combined_time_series.csv")
data$Date <- as.Date(data$Date)

# Core transforms used earlier
log_CPI <- log(data$CPI)
diff_log_cpi <- diff(log_CPI)
final_CPI <- diff_log_cpi

lambda_ppi <- BoxCox.lambda(data$PPI)
PPI_bc <- BoxCox(data$PPI, lambda = lambda_ppi)
final_PPI <- diff(PPI_bc)

housing_bc <- BoxCox(data$HousingIndex, lambda = "auto")
diff_housing <- diff(housing_bc)
final_housing <- diff(diff_housing) # as used previously

final_Unemp <- diff(data$Unemployment)
final_FedFund <- diff(data$FedFunds)

# ===== Align as zoo and merge (keep common intersection) =====
z_CPI <- zoo(final_CPI, order.by = data$Date[-1])
z_PPI <- zoo(final_PPI, order.by = data$Date[-1])
z_House <- zoo(final_housing, order.by = data$Date[-(1:2)])
z_Unemp <- zoo(final_Unemp, order.by = data$Date[-1])
z_Fund <- zoo(final_FedFund, order.by = data$Date[-1])

Z <- merge(CPI = z_CPI,
           PPI = z_PPI,
           FedFund = z_Fund,
           Housing = z_House,
           all = FALSE)

covid_step_all <- zoo(
  as.integer(data$Date >= as.Date("2020-12-01") & data$Date <= as.Date("2022-05-01")),
  order.by = data$Date
)
Z2 <- merge(Z, COVID = covid_step_all, all = FALSE)

# Helper to convert zoo -> monthly ts (preserves calendar start)
zoo_to_ts <- function(z) {
  idx <- index(z)
  start <- c(as.integer(format(min(idx), "%Y")), as.integer(format(min(idx), "%m")))
  ts(coredata(z), start = start, frequency = 12)
}

Y_full <- zoo_to_ts(Z2[, c("CPI", "PPI", "FedFund", "Housing")])
X_full <- zoo_to_ts(Z2[, "COVID", drop = FALSE])

vars_in <- c("PPI", "FedFund", "Housing")
y_model <- Y_full[, c("CPI", vars_in)]
x_model <- X_full

# ===== Force test = Jan-Dec 2024, train = everything strictly before Jan 2024 =====
y_dates <- as.Date(as.yearmon(time(y_model)))

test_start <- as.Date("2024-01-01")
test_end <- as.Date("2024-12-01") # monthly timestamp convention

train <- which(y_dates < test_start)
test <- which(y_dates >= test_start & y_dates <= test_end)

if (length(test) != 12) {
  stop(paste0("Expected 12 test months in 2024, but found ", length(test),
    ". Check your input dates; y_dates range is ", min(y_dates), " to ", max(y_dates),
    "."))
}

# Confirm split dates
last_train_date <- max(y_dates[train]) # should be 2023-12-01
message("Train ends at: ", last_train_date)

## Train ends at: 2023-12-01

message("Test spans: ", min(y_dates[test]), " to ", max(y_dates[test]))

## Test spans: 2024-01-01 to 2024-12-01

# ===== Fit VAR (choose p as before or reselect) =====
p <- 6 # keep your chosen lag
var_mod <- VAR(y_model[train, ], p = p, type = "const", exogen = x_model[train, , drop = FALSE], season = 12)

# ===== Forecast exactly over 2024 =====
h <- length(test)
fc <- predict(var_mod, n.ahead = h, dumvar = x_model[test, , drop = FALSE], ci = 0.95)

# Build ts objects for CPI difflg forecasts and actuals aligned to 2024
cpi_fc <- ts(fc$fcst$CPI[, "fcst"], start = time(y_model)[test][1], frequency = 12)
cpi_act <- y_model[test, "CPI"]

# ===== Invert back to level =====
# Grab last observed log(CPI) right before 2024 (i.e., Dec 2023)
last_log_cpi_train <- log(data$CPI[match(last_train_date, data$Date)])
if (is.na(last_log_cpi_train)) {
  # fallback if dates are not exact day matches
  closest_idx <- which.min(abs(as.numeric(data$Date - last_train_date)))
  last_log_cpi_train <- log(data$CPI[closest_idx])
}

inverted_fcst_cpi <- exp(cumsum(cpi_fc) + last_log_cpi_train)
inverted_actual_cpi <- exp(cumsum(cpi_act) + last_log_cpi_train)

inverted_fcst_cpi

## [1] 308.1562 309.3940 310.7020 311.6589 312.3027 313.0588 313.3015 313.8215
## [9] 314.5826 314.8843 314.5416 314.0634

inverted_actual_cpi

## [1] 308.417 310.326 312.332 313.548 314.069 314.175 314.540 314.796 315.301
## [10] 315.664 315.493 315.605

lower_cpi_test <- as.numeric(fc$fcst$CPI[, "lower"])
upper_cpi_test <- as.numeric(fc$fcst$CPI[, "upper"])

inverted_lower_cpi <- exp(cumsum(lower_cpi_test) + last_log_cpi_train)
inverted_upper_cpi <- exp(cumsum(upper_cpi_test) + last_log_cpi_train)

# ===== Plotting data frames =====
fc_df <- tibble(
  Date = y_dates[test],
  Forecast = as.numeric(inverted_fcst_cpi),
  Lower = as.numeric(inverted_lower_cpi),
  Upper = as.numeric(inverted_upper_cpi)
)

actual_df <- tibble(
  Date = data$Date,
  Actual = data$CPI
)

rmse_inv <- sqrt(mean((inverted_fcst_cpi - inverted_actual_cpi)^2, na.rm = TRUE))
message("Test RMSE (original scale, 2024): ", round(rmse_inv, 3))

## Test RMSE (original scale, 2024): 1.239

# ===== Deck-ready ggplot objects =====
base_theme <- theme_minimal(base_size = 13) +
  theme(
    panel.grid.minor = element_blank(),
    panel.grid.major.x = element_line(linewidth = 0.25),
    panel.grid.major.y = element_line(linewidth = 0.25),
    plot.title = element_text(face = "bold"),
    plot.subtitle = element_text(margin = margin(b = 8)),
    axis.title.x = element_blank(),
    legend.position = "top",
    legend.title = element_blank()
  )

# 1) Full history with clear split at Dec 2023 (kept)
p_full <- actual_df %>%
  ggplot(aes(x = Date)) +
  geom_line(aes(y = Actual, color = "Actual")) +
  geom_vline(xintercept = as.numeric(last_train_date), linetype = "dashed") +
  geom_ribbon(data = fc_df, aes(x = Date, ymin = Lower, ymax = Upper, fill = "95% PI"), alpha = 0.15, inherit.aes = FALSE) +
  geom_line(data = fc_df, aes(y = Forecast, color = "Forecast"), linewidth = 0.9) +
  scale_x_date(labels = label_date("%Y")) +
  scale_color_manual(values = c("Actual" = "#1f1f1f", "Forecast" = "#d62728")) +
  scale_fill_manual(values = c("95% PI" = "#9ecae1")) +
  labs(title = "CPI: Inverted VAR Forecast vs Actual",
    subtitle = paste0("Train/Test split at ", format(last_train_date, "%b %Y"),
      "\nTest RMSE (orig scale, 2024): ", round(rmse_inv, 2)),
    y = "CPI (Index)") +
  base_theme

# 2) Zoomed 2024 view WITHOUT the split line
zoom_start <- as.Date("2023-01-01")
zoom_end <- as.Date("2024-12-31")


p_zoom <- actual_df %>%
  filter(Date >= zoom_start & Date <= zoom_end) %>%
  ggplot(aes(x = Date)) +
  geom_line(aes(y = Actual, color = "Actual")) +
  # intentionally removed geom_vline here
  geom_ribbon(data = fc_df, aes(x = Date, ymin = Lower, ymax = Upper, fill = "95% PI"), alpha = 0.15, inherit.aes = FALSE) +
  geom_line(data = fc_df, aes(y = Forecast, color = "Forecast"), linewidth = 0.9) +
  scale_x_date(date_breaks = "1 month", date_labels = "%b %Y", limits = c(zoom_start, zoom_end)) +
  scale_color_manual(values = c("Actual" = "#1f1f1f", "Forecast" = "#d62728")) +
  scale_fill_manual(values = c("95% PI" = "#9ecae1")) +
  labs(title = "CPI: 2024 Forecast Horizon",
    subtitle = "Full Jan-Dec 2024 test set",
    y = "CPI (Index)") +
  base_theme +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# 3) Residuals (training period)
train_dates_aligned <- y_dates[train]
res_vec <- as.numeric(resid(var_mod)[, "CPI"]) # length = length(train) - p
res_dates <- tail(train_dates_aligned, length(res_vec))
resid_df <- tibble(Date = res_dates, Resid = res_vec)

p_resid <- ggplot(resid_df, aes(Date, Resid)) +
  geom_hline(yintercept = 0, linewidth = 0.3, color = "#666666") +
  geom_line(color = "#3182bd") +
  labs(title = "VAR(CPI) Residuals (Training)", y = "Residual") +
  base_theme

# ===== Print / Save =====
print(p_full)
```

CPI: Inverted VAR Forecast vs Actual
Train/Test split at Dec 2023
Test RMSE (orig scale, 2024): 1.24



CPI: 2024 Forecast Horizon
Full Jan-Dec 2024 test set

